

EECS 498 · AASE · FALL 2026

Welcome to AASE

LAB 00 · SETUP LAB · AUG 31 – SEP 1, 2026

STEP 1 • 10 MINUTES • START NOW

Join Ed, take the survey

1. Open Canvas. Click the **Ed** link. Join
2. Find the survey post on Ed. Open it



Three pieces, two machines

```
Ollama    serves a model, on a port  
.env      one line: which port that is  
aider     reads .env, talks to whatever answers there
```

- **Steps 2–4:** all three on your laptop
- **Step 5:** all three again on a CAEN machine, which has a GPU

The steps

WORK AT YOUR OWN PACE. RAISE A HAND.

Clone `aider-ollama`

```
git clone https://github.com/eecs498-aase/aider-ollama.git
cd aider-ollama
python3 -m pip install aider-install && aider-install
```

- Its **README** has every step below, in more detail

Serve a model

1. Install Ollama from ollama.com/download
2. Then, in the `aider-ollama` clone:

```
ollama serve &  
ollama pull qwen3.5:4b
```

About 3 GB. Let it finish.

Point aider at it

```
cp .env.local .env  
bin/check  
aider
```

All green? You're done with the part that matters. A red line? Read it, then grab a TA.

Same thing, with a GPU

```
bin/install-aider  
bin/install-ollama  
  
loginctl enable-linger $USER  
ollama serve & ollama pull qwen3.5:9b  
cp .env.local .env && bin/check
```

Pull once. Your models are there at every lab machine.

BEFORE YOU LEAVE

Checklist

- On Ed
- Survey submitted
- `bin/check` all green

Anything red? Read the repo's README, then grab a TA. There's no lab next week.

Questions?