EECS 498 · AASE · FALL 2026

Welcome to AASE

LAB 00 · SETUP LAB · AUG 31 – SEP 1, 2026

STEP 1 · 10 MINUTES · START NOW

Join Ed, take the survey

1. Open Canvas. Click the **Ed** link. Join
2. Find the survey post on Ed. Open it

Local LLMs on your machine

- `ollama pull` downloads a model. `ollama serve` runs it
- Free, private, no API key. It's just *yours*
- Small models fit a laptop. Bigger ones want the CAEN GPUs

Our programming tool

- A pair programmer in your terminal: describe a change, it edits your files
- Every edit is a git commit, so you can always *undo*
- It talks to whatever model Ollama is running. That’s the whole stack

The steps

WORK AT YOUR OWN PACE. RAISE A HAND.

STEP 2 · ON YOUR LAPTOP

Clone `aider-ollama`

```
git clone https://github.com/eecs498-aase/aider-ollama.git
cd aider-ollama
python3 -m pip install aider-install && aider-install
```

- Its **README** has every step below, in more detail

STEP 3 · ON YOUR LAPTOP

Serve a model

1. Install Ollama from **ollama.com/download**
2. Then, in the `aider-ollama` clone:

```
ollama serve &
ollama pull qwen3.5:4b
```

About 3 GB. Let it finish.

STEP 4 · ON YOUR LAPTOP

Point aider at it

```
cp .env.local .env
bin/check
aider
```

All green? You're done with the part that matters. A red line? Read it, then grab a TA.

STEP 5 · ON A CAEN MACHINE

Same thing, with a GPU

```
bin/install-aider
bin/install-ollama

loginctl enable-linger $USER
ollama serve & ollama pull qwen3.5:9b
cp .env.local .env && bin/check
```

Pull once. Your models are there at every lab machine.

BEFORE YOU LEAVE

Checklist

- On Ed
- Survey submitted
- `bin/check` all green

Anything red? Read the repo's README, then grab a TA. There's no lab next week.

Questions?