

EECS 498 · AASE · FALL 2026

Modern agentic CLIs

LECTURE 05 · SEP 15, 2026

WHERE WE ARE

Lab01 left one question open

You followed one request through six steps: assemble, answer, parse, apply, repair, commit.


At every step, the thing that picked what happened next was you, or a rule somebody wrote down in advance.

Who picks the **next action**?

The autonomy gradient

FIVE RUNGS, AND WHERE YOU HAVE BEEN STANDING

Five rungs, not a ranking



The blue rung is where you have spent two weeks.

You are on rung 3

At rung 3

You hand over	One bounded edit task
You choose	The files, and where the task ends
It automates	Matching, repair, tests, the commit
You check	Between tasks

Real automation inside a boundary you drew.

Higher rungs spend more to ask less



Cost per unit of work

A run that searches, reads five files and runs your tests three times spends far more attention than one bounded edit.



Decisions you make yourself

Fewer choices per finished outcome, and less of your attention spent on each one.

Rung 3 fits a change you can name

A two-line fix in a file you can point at is cheaper and safer at rung 3.

Paying rung 4 prices for it buys nothing.

The skill is picking the rung, and you cannot pick what you cannot see.

What makes a tool agentic

FOUR PRIMITIVES, AND A TEST YOU CAN APPLY

Four primitives, not a vibe



Loop

Takes another step
without being asked.



Tools

Actions beyond
emitting text.



Planning

Breaks a goal into
steps.



Memory

Something survives
the turn.

"Agentic" is used to mean "recent." These four are checkable.

Aider has all four

Primitive	Where it already is in Aider
Loop	A failed edit becomes the next prompt; tests can feed back
Tools	Reads, writes, commits, runs the commands you configured
Planning	A plan drafted in prose, stopping for a human to read
Memory	Session history, plus a repo map of code you never added

The difference is **who chooses.**

THE MOVE

One choice changes the rung

At rung 3 you choose the files, the boundary, and whether to run the plan.

Move the choosing of the next action from the harness to the model, and everything people find surprising about agentic tools follows.

One change. Not a new category of intelligence.

Chat break

FIVE MINUTES, ON THE DESIGN YOU DRAFTED YESTERDAY

Move your own design up a rung

Your assistant sits at rung 3, and your specification puts it there.

Pick **one** primitive you would add: the model choosing its next action, discovering files you never selected, or carrying state between sessions.

Then say what you would have to build first to make it safe.

PollEv

<https://PollEv.com/marcusdarden365>

The same request, one rung up

LIVE, ON THIS MACHINE

You know what this took at rung 3

Add a `--priority` flag to `taskr`'s `add` command.

You added three files, because a file you do not add is not in the context. You framed the task. You read the diff. You found the missing write yourself.

Nobody types /add

It searches

Finds where arguments are parsed, without being told.

It reads

Follows the code to the files the change actually touches.

It runs

Executes the tests and reacts to what they say.

Context assembly became something the program does by acting.

Each result chooses the next step

```
search "add_argument" → three hits, one in cli.py
read cli.py           → finds the parser
read store.py         → finds the save path
edit cli.py, task.py, store.py
run tests              → 1 failed
edit store.py
run tests              → passed
```

Nothing in that order was written down in advance, by anyone.

HONESTY

Today's run is a fact about today


What you just watched is what this model and this build did with this repository, this afternoon.

A run that goes badly is more useful than one that goes well. Do not rescue it too quickly.

What the harness still decides

THE PART THE DEMO HIDES

The model picks from a menu it did not write



Somebody decided a search tool exists and a database client does not.

Permissions are the real safety layer

You already built this shape. Nothing changes on disk until a human approves it, and that holds whatever the model says, because your program enforces it.

The same question scales up: when the model can run commands, which ones run without being asked?

Prose asks. Code and permissions
enforce.

A loop needs a reason to stop

A goal it can recognize as met. A budget it can spend. A human who interrupts.

At rung 3 the task boundary was the stop condition, and you drew it by hand every single time.

YOUR BUILD

Rung 3 is the point, not a limitation

Exact matching. One approval before any write. One commit per accepted edit set. A guarded undo.

Those are what make the state after a failure something you can say in a sentence and assert in a test.

Build that, and you have built what a loop needs underneath it.

The trace is the evidence

When a run produces the wrong outcome, the output does not tell you why.

Which step picked the wrong file? Which result did it misread? Which missing tool did it improvise around?

"The model is bad at this" is the answer people reach for, and rarely the one the evidence supports.

THE WEEK

The build is due in three weeks

When	What
Thu Sep 17	L06 closes Apply
Mon 21 / Tue 22	Lab02, supported build time
Thu Sep 24, evening	Hackathon 1, on your own project
Tue Oct 6, 11:59 PM	The build is due

TAKE THESE WITH YOU

Three questions

1. Where on the gradient is the tool you reach for by default, and did you choose that or inherit it?
2. What would have to be true before you would let a model run a command without asking you?
3. What would you need to see in a trace to be convinced a run failed for a reason other than the model?

Thursday, the loop stops being
someone else's.

Questions?