

EECS 498 · AASE · FALL 2026

Hackathon 1 briefing

LECTURE 08 · SEP 24, 2026

CHAT BREAK

Habits with a stronger model

Tonight Aider runs a far stronger model, still on a 32K budget.

Write one line: one habit from aider-practice that matters *more* with a stronger model, and one that matters less.

Be ready to say why.

TONIGHT

The evening at a glance

When	Tonight, 7 to 9 PM
Where	Leinweber 1355
Work	Individual
Weight	9% of your grade
Bring	Your laptop and its charger

Tonight is about **taskr**. The rest is
in your repository at **seven**.

How the evening works

- **Modular:** do as much as you can do well
- **Judged on evidence**, not feature count
- **The brief** is in your repository
- **Submit** by pushing before 9 PM

The model

GPT-5.6-LUNA, ON YOUR OWN KEY

From qwen to gpt-5.6-luna

	So far	Tonight
Model	qwen3.5:9b	gpt-5.6-luna
Runs on	Your laptop	OpenAI, through U-M
Context	32K tokens	32K, by design
Class	Small, open weights	Frontier, hosted
Speed	Your hardware	About 80 tokens a second

A **strong** model, on the **32K** budget
you know.

The Toolkit key and its budget

- **Where:** toolkit.umgpt.umich.edu
- **Sign in:** your U-M account
- **Credit:** \$25, the whole budget
- **Model:** `gpt-5.6-luna` only

GIVEN: TASKR/MODEL.PY

One request to the model

```
body = {"model": "gpt-5.6-luna", "messages": messages}
if tools:
    body["tools"] = tools                # actions the model may ask for
headers = {
    "Content-Type": "application/json",
    "Authorization": f"Bearer {api_key}", # your Toolkit key
}
request = urllib.request.Request(
    base_url + "/chat/completions",      # base_url ends in /v1
    data=json.dumps(body).encode(), headers=headers)
with urllib.request.urlopen(request) as response:
    reply = json.load(response)
message = reply["choices"][0]["message"] # content, or tool_calls
```

You call it: `call(messages, base_url=..., model=..., api_key=..., tools=...)`

What \$25 buys

Run	Cost
A scripted run of the whole evening	\$0.24
Typical evening, estimated	About \$1
Heavy evening, estimated	\$2 to \$4

The way to spend all \$25: calling the model over and over, with nothing to stop it.

Using a strong model well

- **The 32K budget:** `/drop` files, `/clear` between tasks
- **Small commits:** one working step each
- **"Add file to the chat?":** usually N

Setup and recording

BEFORE YOU ARRIVE, AND THE FIRST TEN MINUTES

BEFORE TONIGHT

Before you arrive

- **Laptop** and charger
- **Python 3.10+**, git, Aider 0.86.2
- **Your Toolkit key**, saved
- **GitHub push** without a password prompt

The first ten minutes

1. **Check in** on the form
2. **Accept** your GitHub invitation
3. **Clone** your repository
4. **Run** the setup script

The exact steps are on the screen tonight.

REQUIRED

Everyone runs bin/setup

1. Run `bin/setup`
2. **Paste** your Toolkit key when it asks
3. **Check** that every step says ok

Anything says PROBLEM? Tell us straight away.

What gets recorded

- **Every Aider commit** carries its conversation
- **Every commit** snapshotted to `safety/<machine>`
- `/undo` works, and the undone change is kept
- `bin/aider-log` lists it all

SETUP CHECK, NOW

The setup check

```
python3 --version      # 3.10 or newer  
git --version  
aider --version        # 0.86.2  
ssh -T git@github.com  # greets you by name
```

Then sign in at toolkit.umgpt.umich.edu and check your key exists.

Questions?